

# Exercise 12 – Bayesian inference and Data assimilation

Group : Duckies

## 1.1

The variance of the sample mean  $\bar{g}_X$  can be calculated as follows:

Since  $X_i$  are iid samples from the distribution  $\pi$ ,  $g(X_i)$  are also iid. Let  $\sigma^2 = \text{Var}(g(X))$ . Then, by the properties of variance, we have:

$$\text{Var}(\bar{g}_X) = \text{Var}\left(\frac{1}{N} \sum_{i=1}^N g(X_i)\right) = \frac{1}{N^2} \text{Var}\left(\sum_{i=1}^N g(X_i)\right) = \frac{1}{N^2} \sum_{i=1}^N \text{Var}(g(X_i)) = \frac{N\sigma^2}{N^2} = \frac{\sigma^2}{N}$$

So, the variance of the sample mean  $\bar{g}_X$  is  $\frac{\sigma^2}{N}$ , where  $\sigma^2$  is the variance of  $g(X)$ .

## 1.2

This is a result of the importance sampling technique. The idea is to estimate the expected value of  $g(X)$  with respect to the distribution  $\pi$  by using samples from another distribution  $\pi'$ .

We can derive the result as follows:

$$\pi[g] = \mathbb{E}_\pi[g(X)] = \int g(x)\pi(x)dx = \int g(x)\frac{\pi(x)}{\pi'(x)}\pi'(x)dx = \mathbb{E}_{\pi'}\left[\frac{\pi(Y)}{\pi'(Y)}g(Y)\right]$$

So, we have shown that  $\pi[g] = \mathbb{E}\left[\frac{\pi(Y)}{\pi'(Y)}g(Y)\right]$ .

## 1.3

The variance of the new sample mean  $\bar{g}_Y$  can be calculated as follows:

Let  $Z_i = \frac{\pi(Y_i)}{\pi'(Y_i)}g(Y_i)$ . Since  $Y_i$  are iid samples from the distribution  $\pi'$ ,  $Z_i$  are also iid. Let  $\sigma_Z^2 = \text{Var}(Z_i)$ . Then, by the properties of variance, we have:

$$\text{Var}(\bar{g}_Y) = \text{Var}\left(\frac{1}{N} \sum_{i=1}^N Z_i\right) = \frac{1}{N^2} \text{Var}\left(\sum_{i=1}^N Z_i\right) = \frac{1}{N^2} \sum_{i=1}^N \text{Var}(Z_i) = \frac{N\sigma_Z^2}{N^2} = \frac{\sigma_Z^2}{N}$$

So, the variance of the new sample mean  $\bar{g}_Y$  is  $\frac{\sigma_Z^2}{N}$ , where  $\sigma_Z^2$  is the variance of  $Z_i$ .

## 1.4

Let's denote  $h(Y) = \frac{\pi(Y)}{\pi'(Y)}g(Y)$ . Then, the variance of  $\bar{g}_Y$  can be expressed as:

$$\text{Var}(\bar{g}_Y) = \frac{1}{N} \text{Var}_{\pi'}(h(Y)) = \frac{1}{N} (\mathbb{E}_{\pi'}[h(Y)^2] - \mathbb{E}_{\pi'}[h(Y)]^2)$$

Since  $\mathbb{E}_{\pi'}[h(Y)] = \pi[g]$  is a constant, we can minimize the variance of  $\bar{g}_Y$  by minimizing  $\mathbb{E}_{\pi'}[h(Y)^2]$ .

Let's expand the expression for  $\mathbb{E}_{\pi'}[h(Y)^2]$ :

$$\mathbb{E}_{\pi'}[h(Y)^2] = \int h(y)^2 \pi'(y) dy = \int \left( \frac{\pi(y)}{\pi'(y)} g(y) \right)^2 \pi'(y) dy = \int \frac{\pi(y)^2}{\pi'(y)} g(y)^2 dy$$

To minimize this expression, we can use the Cauchy-Schwarz inequality:

$$\int \frac{\pi(y)^2}{\pi'(y)} g(y)^2 dy \geq \left( \int \sqrt{\frac{\pi(y)^2}{\pi'(y)}} |g(y)| \sqrt{\pi'(y)} dy \right)^2 = \left( \int \pi(y) |g(y)| dy \right)^2$$

The equality holds when  $\sqrt{\frac{\pi(y)^2}{\pi'(y)}} |g(y)| = k \sqrt{\pi'(y)}$  for some constant  $k$ . Solving for  $\pi'(y)$ , we get:

$$\begin{aligned} \sqrt{\frac{\pi(y)^2}{\pi'(y)}} |g(y)| &= k \sqrt{\pi'(y)} \\ \Leftrightarrow \frac{\pi(y)^2}{\pi'(y)} g(y)^2 &= k^2 \pi'(y) \\ \Leftrightarrow \pi'(y) &= \frac{\pi(y)^2 g(y)^2}{k^2} \end{aligned}$$

Since  $\int \pi'(y) dy = 1$ , we have:

$$k^2 = \int \frac{k^2 \pi(y)^2 g(y)^2}{k^2} dy = \int \pi(y)^2 g(y)^2 dy$$

So, the optimal choice of  $\pi'$  that minimizes the variance of  $\bar{g}_Y$  is:

$$\boxed{\pi'(y) = \frac{\pi(y)^2 g(y)^2}{\int \pi(x)^2 g(x)^2 dx}}$$

In practice, however, it might not be possible to achieve this optimal choice of  $\pi'$ , and one would need to propose a suitable distribution that can represent  $\pi[g]$  well or is easy to sample from.

## 1.5

Here is a pseudo-code for the importance resampling algorithm to generate  $N$  samples from  $\pi$ :

1. Generate  $N$  samples  $Y_i$  from the proposal distribution  $\pi'$
2. Calculate the unnormalized weights  $\tilde{w}_i = C * \pi(Y_i) / \pi'(Y_i)$   
for  $i = 1, \dots, N$
3. Normalize the weights  $w_i = \tilde{w}_i / \text{sum}(\tilde{w}_i)$  for  $i = 1, \dots, N$

4. Resample  $X_i$  from  $Y_i$  with probability mass function  $w_i$  for  $i = 1, \dots, N$
5. Return the resampled values  $X_i$  for  $i = 1, \dots, N$

This algorithm generates  $N$  samples  $Y_i$  from the proposal distribution  $\pi'$ , calculates the importance weights  $w_i$ , and then resamples  $X_i$  from  $Y_i$  with probability mass function  $w_i$ . The resulting samples  $X_i$  can be used to approximate the distribution  $\pi$ .

## 2.1

Here is a pseudo-code for implementing the particle filter algorithm:

1. Initialize  $N$  particles  $x_0^i$  from the prior distribution  $p_0$  for  $i = 1, \dots, N$
2. Set the initial weights  $w_0^i = 1/N$  for  $i = 1, \dots, N$
3. for  $n = 1$  to  $T$  do
  4. for  $i = 1$  to  $N$  do
    5. Sample  $x_n^i$  from the transition kernel  $p_n(x_n | x_{n-1}^i)$
    6. Update the weight  $w_n^i = w_{n-1}^i * q_n(y_n | x_n^i)$
  7. end for
  8. Normalize the weights  $w_n^i = w_n^i / \sum(w_n^i)$  for  $i = 1, \dots, N$
  9. Resample  $(x_n^i, w_n^i)$  with probability mass function  $w_n^i$  for  $i = 1, \dots, N$
  10. Set the weights of the resampled particles to  $w_n^i = 1/N$  for  $i = 1, \dots, N$
  11. end for
12. Return  $(x_T^i, w_T^i)$  for  $i = 1, \dots, N$

This algorithm generates  $N$  particles  $x_0^i$  from the prior distribution  $p_0$  and sets their initial weights to  $w_0^i = 1/N$ . Then, it iterates over time from  $n = 1$  to  $T$ . At each time step, it samples new particles  $x_n^i$  from the transition kernel  $p_n(x_n | x_{n-1}^i)$  and updates their weights using the observation likelihood  $q_n(y_n | x_n^i)$ . The weights are then normalized and the particles are resampled with replacement according to their weights. The weights of the resampled particles are reset to  $w_n^i = 1/N$ . The final output is the set of particles and their weights at time  $T$ , which approximates the posterior distribution  $\pi_{X_T | Y_{1:T}}$ .

## 2.2

Here is an updated version of the pseudo-code that includes a resampling step based on the effective sample size (ESS):

1. Initialize  $N$  particles  $x_0^i$  from the prior distribution  $p_0$  for  $i = 1, \dots, N$

2. Set the initial weights  $w_{\theta^i} = 1/N$  for  $i = 1, \dots, N$
3. for  $n = 1$  to  $T$  do
4.     for  $i = 1$  to  $N$  do
5.         Sample  $x_n^i$  from the transition kernel  $p_n(x_n \mid x_{n-1}^i)$
6.         Update the weight  $w_n^i = w_{n-1}^i * q_n(y_n \mid x_n^i)$
7.     end for
8.     Normalize the weights  $w_n^i = w_n^i / \text{sum}(w_n^i)$  for  $i = 1, \dots, N$
9.     Calculate the effective sample size  $ESS = 1 / \text{sum}(w_n^i)^2$
10.     if  $ESS < \gamma * N$  then
11.         Resample  $(x_n^i, w_n^i)$  with probability mass function  $w_n^i$
- for  $i = 1, \dots, N$
12.         Set the weights of the resampled particles to  $w_n^i = 1/N$
- for  $i = 1, \dots, N$
13.     end if
14. end for
15. Return  $(x_T^i, w_T^i)$  for  $i = 1, \dots, N$

This updated version of the algorithm includes an additional input parameter  $\gamma \in (0, 1)$  that determines the threshold for resampling based on the effective sample size (ESS). At each time step, after updating and normalizing the weights, the algorithm calculates the ESS as  $ESS = 1 / \sum (w_n^i)^2$ . If the ESS is less than  $\gamma N$ , then the particles are resampled with replacement according to their weights and their weights are reset to  $w_n^i = 1/N$ . This resampling step helps to prevent weight degeneracy and maintain a large effective sample size.

## 3.1

Here is a Python code that implements the importance resampling algorithm to generate 10,000 samples of  $X$  from the given mixture of Gaussian random variables using a Gaussian proposal distribution with  $\sigma^2 = 1$  or  $\sigma^2 = 4$ :

```
In [2]: import numpy as np
from scipy.stats import norm

def importance_resampling(N, sigma2):
    # Set the proposal distribution pi' = N(theta, sigma^2)
    pi_prime = norm(loc=0, scale=np.sqrt(sigma2))

    # Generate N samples Y_i from the proposal distribution pi'
    Y = pi_prime.rvs(size=N)

    # Define the target distribution pi(x)
    def pi(x):
```

```

        return 0.5 * norm.pdf(x, loc=-4, scale=1) + 0.5 * norm.pdf(x, loc=4,
scale=1)

    # Calculate the unnormalized weights  $\tilde{w}_i = \pi(Y_i) / \pi'(Y_i)$ 
    tilde_w = pi(Y) / pi_prime.pdf(Y)

    # Normalize the weights  $w_i = \tilde{w}_i / \sum(\tilde{w}_i)$ 
    w = tilde_w / np.sum(tilde_w)

    # Resample  $X_i$  from  $Y_i$  with probability mass function  $w_i$ 
    X = np.random.choice(Y, size=N, p=w)

    return X

# Generate 10,000 samples of X using a Gaussian proposal with  $\sigma^2 = 1$ 
X1 = importance_resampling(N=10000, sigma2=1)

# Generate 10,000 samples of X using a Gaussian proposal with  $\sigma^2 = 4$ 
X2 = importance_resampling(N=10000, sigma2=4)

```

This code defines a function `importance_resampling` that takes as input the number of samples `N` and the variance `sigma2` of the Gaussian proposal distribution. The function generates `N` samples `Y` from the Gaussian proposal distribution `pi_prime`, calculates the importance weights `w`, and then resamples `X` from `Y` with probability mass function `w`. The resulting samples `X` can be used to approximate the target distribution `pi(x)`. The code generates two sets of samples `X1` and `X2` using different proposal distributions with `sigma2 = 1` and `sigma2 = 4`, respectively.

## 3.2

Here is a Python code that plots the histograms of the samples generated using the importance resampling algorithm with different proposal distributions, and overlays them with the true density  $\pi(x)$ :

In [3]:

```

import matplotlib.pyplot as plt

# Define the target distribution  $\pi(x)$ 
def pi(x):
    return 0.5 * norm.pdf(x, loc=-4, scale=1) + 0.5 * norm.pdf(x, loc=4,
scale=1)

# Plot the histogram of the samples generated using a Gaussian proposal with
 $\sigma^2 = 1$ 

```


```
plt.hist(X1, bins=30, density=True, alpha=0.5, label='sigma^2 = 1')

# Plot the histogram of the samples generated using a Gaussian proposal with
sigma^2 = 4
plt.hist(X2, bins=30, density=True, alpha=0.5, label='sigma^2 = 4')

# Overlay the true density pi(x)
x = np.linspace(-10, 10, 1000)
plt.plot(x, pi(x), 'k', label='pi(x)')

# Add Labels and Legend
plt.xlabel('x')
plt.ylabel('Density')
plt.legend()

# Show the plot
plt.show()
```



This code uses the `matplotlib` library to plot the histograms of the samples `X1` and `X2` generated using different proposal distributions with `sigma2 = 1` and `sigma2 = 4`, respectively. It also plots the true density `pi(x)` as a black line on top of the histograms.

The choice of `sigma2` for the proposal distribution can affect the quality of the samples generated by the importance resampling algorithm. If `sigma2` is too small, then the proposal distribution might not cover the entire support of the target distribution `pi(x)`, resulting in samples that are concentrated in a small region of the space. On the other hand, if `sigma2` is too large, then the proposal distribution might generate many samples in regions where `pi(x)` is very small, resulting in a large variance for the importance weights and a high variance for the importance sampling estimator. A good choice of `sigma2` should balance these two considerations and generate samples that are representative of the target distribution `pi(x)`.